**Methodological Note: for Finnegans Jape

Polyphonic Conditioning of a Small Language Model on *Finnegans Wake***

Overview

This project develops a small language model (SLM) conditioned on *Finnegans Wake* through a hybrid textual–prosodic methodology. Rather than treating the Wake as a static text to be memorized or summarized, the model is trained to **inhabit the text's rhythmic, polyphonic, and interpretive affordances**. The approach combines light parameter-efficient fine-tuning with retrieval-based contextualization and audio-informed temporal conditioning.

The guiding assumption is that *Finnegans Wake* is not primarily a semantic object but a **performed linguistic system**, historically and socially realized through voice, cadence, overlap, mishearing, and collective reading.

Core Design Principles

1. **Polyphony over singular authority**
   The Wake resists definitive voice or interpretation. The model is therefore designed to reflect lawful variation rather than canonical reading.

2. **Rhythm as structural signal**
   Temporal features—pause, overlap, repetition, and breath—are treated as primary informational channels, not decorative additions.

3. **Retrieval over memorization**
   Given the Wake's limited corpus size and extreme density, interpretive flexibility is preserved by retrieval-augmented generation rather than full internalization.

Textual Corpus

The primary textual corpus includes:

- *Finnegans Wake* (public domain text)
- Selected ALP (Anna Livia Plurabelle) passages

- Humanistic glosses (e.g., ALP annotations, HCE marginalia)
- Source texts for the Wake - e.g. Taliesen, Finn McCool, Vico, Ulysses

- Variant editions and editorial notes

Rather than training on the full text indiscriminately, **curated passages** are selected for rhythmic richness and interpretive density.

Audio Conditioning Sources

Two distinct audio sources are used, each serving a different methodological function:

1. Authorial Reading (James Joyce)

Existing recordings of Joyce reading from *Finnegans Wake* are used as **prosodic ground truth**. These recordings provide:

- baseline tempo

- stress and vowel elongation patterns

- tolerance for ambiguity and slippage

- historically situated performance norms

Joyce's recordings are not treated as authoritative interpretations, but as **anchoring constraints** on rhythm and pacing.

2. Contemporary Choral Readings

Group (choral) readings of the same Wake passages are recorded with deliberate overlap, echo, and variation. These recordings encode:

- distributed cognition

- simultaneity and interference

- misreading as productive signal

- communal breath and tempo variation

Choral readings are treated as models of **living Wake performance**, demonstrating how meaning arises socially rather than univocally.

Technical Integration

Audio is not used for speech synthesis or voice cloning. Instead, it functions as **temporal conditioning metadata**, informing:

- generation pacing
- line-breaking and spacing
- thunderword emission timing
- tolerance for ambiguity and overlap

Text and audio are aligned at a coarse-grain level (sectional correspondence rather than word-level transcription), preserving interpretive looseness.

The model learns not to “sound like” Joyce or any chorus, but to **generate text that moves with similar temporal intelligence**.

Model Architecture

- Base model: 7B-class open-source transformer (e.g., Mistral-7B)
- Fine-tuning: LoRA / QLoRA
- Retrieval layer: vector database containing glosses and annotations
- Tokenization: custom tokenizer trained on Wake-specific morphology to preserve portmanteaux and thunderwords

---

Modes of Use

The system is accessed through interpretive modes rather than general chat:

- *murmur* (associative, low-coherence)
- *river* (continuous flow with retrieval)
- *chorus* (polyphonic generation)
- *marginalia* (annotation-aware commentary)

These modes reinforce the model's role as **reading companion**, not explanatory authority.

Epistemic Position

This methodology assumes that learning *Finnegans Wake*—for humans or machines—is not the acquisition of meaning but the **recovery of capacity**: the capacity to move with ambiguity, rhythm, and layered reference.

In this sense, the model functions not as an interpreter of the Wake, but as a **participant in its ongoing performance tradition**.

Conclusion

By combining authorial prosody, communal choral reading, retrieval-based interpretation, and lightweight fine-tuning, this approach offers a scalable, ethical, and philologically sensitive method for training language models on complex literary texts whose meaning resides as much in performance as in inscription.